

Data Warehousing and OLAP: MOLAP and ROLAP

dr. Toon Calders
t.calders@tue.nl

Previous Lectures

- Online analytical processing
 - Data cubes as a conceptual model
 - Query languages for data cubes
- Database explosion problem
 - Materializing the complete cube is impossible
 - Partially materializing can help

This Lecture

- How is the data stored?
 - relational database (ROLAP)
 - Specialized structures (MOLAP)
- How can we speed up computation?
 - Indexing structures
 - bitmap index
 - join index

Implementation

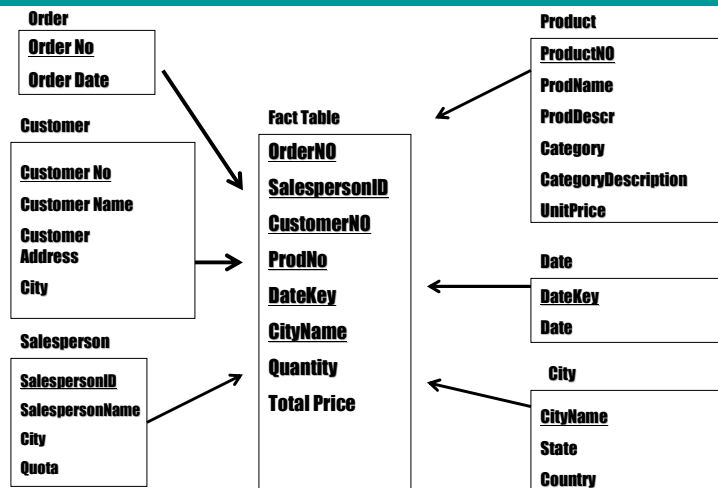
Nowadays systems can be divided in three categories:

- ROLAP (Relational OLAP)
 - OLAP supported on top of a relational database
- MOLAP (Multi-Dimensional OLAP)
 - Use of special multi-dimensional data structures
- HOLAP: (Hybrid)
 - combination of previous two

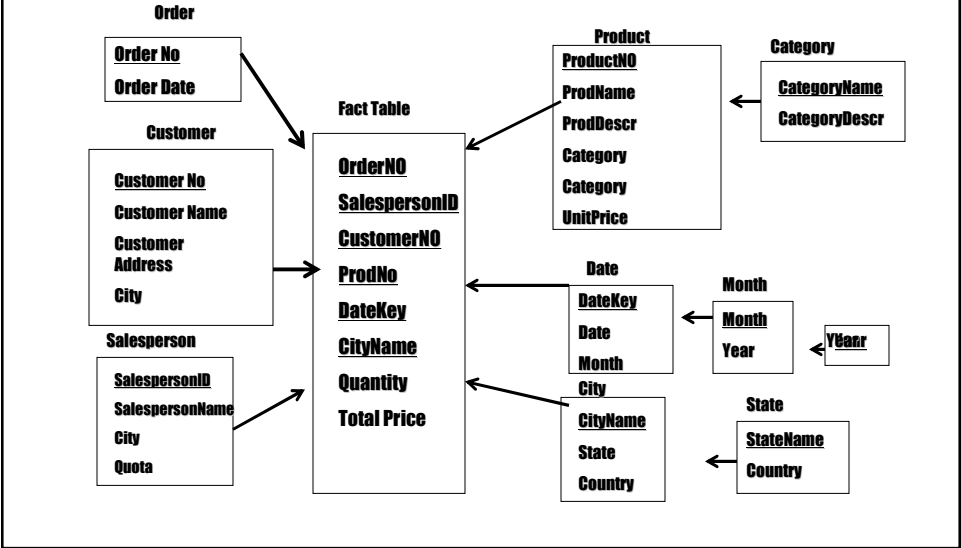
ROLAP

- Typical database scheme:
 - star schema
 - fact table is central
 - links to dimensional tables
 - Extensions:
 - snowflake schema
 - dimensions have hierarchy/extra information attached
 - Star constellation
 - multiple star schemas sharing dimensions

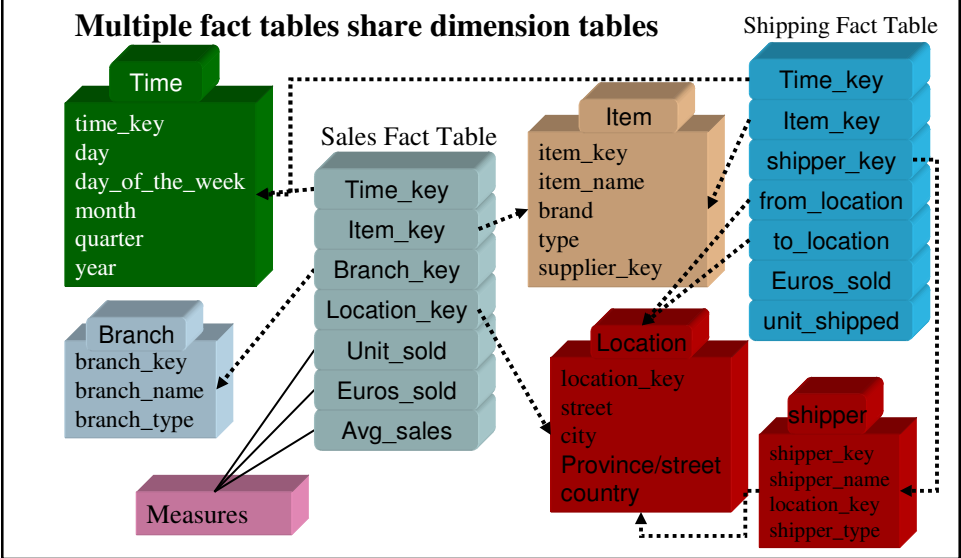
Example of a Star Schema



Example of a Snowflake Schema



Example of Fact Constellation



This Lecture

- How is the data stored?
 - Relational database (ROLAP)
 - **Specialized structures (MOLAP)**
- How can we speed up computation?
 - Indexing structures
 - bitmap index
 - join index

MOLAP


- Not on top of relational database
 - most popular design
 - specialized data structures
 - Multicubes vs Hypercubes
 - Not all subcubes are materialized

Storing the cube

- User identifies set of *sparse* attributes S, and a set of *dense* attributes D.
- Index tree is constructed on sparse dimensions.
- Each leaf points to a multidimensional array indexed by D.

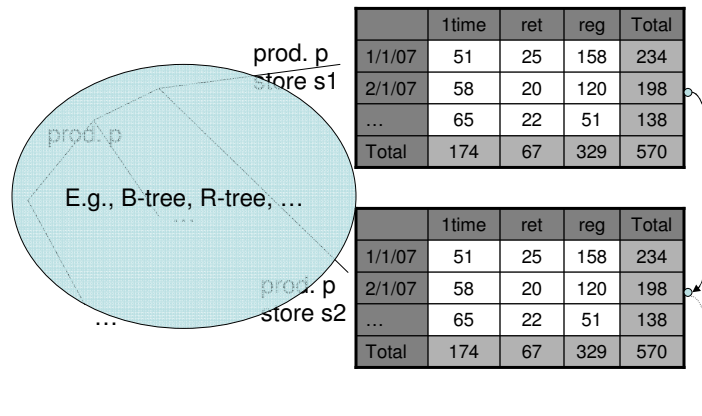
Example

- product, store are sparse dimensions
- date and customer-type are dense



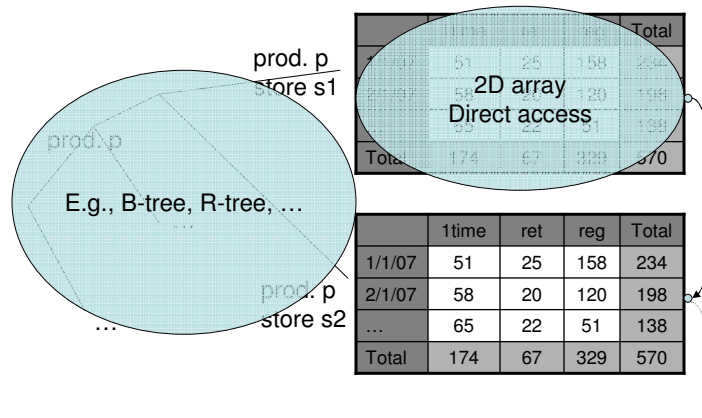
Example

- product, store are sparse dimensions
- date and customer-type are dense



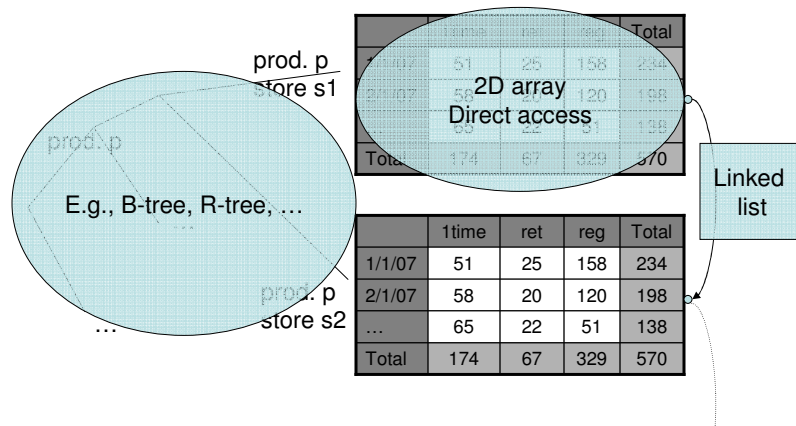
Example

- product, store are sparse dimensions
- date and customer-type are dense



Example

- product, store are sparse dimensions
- date and customer-type are dense

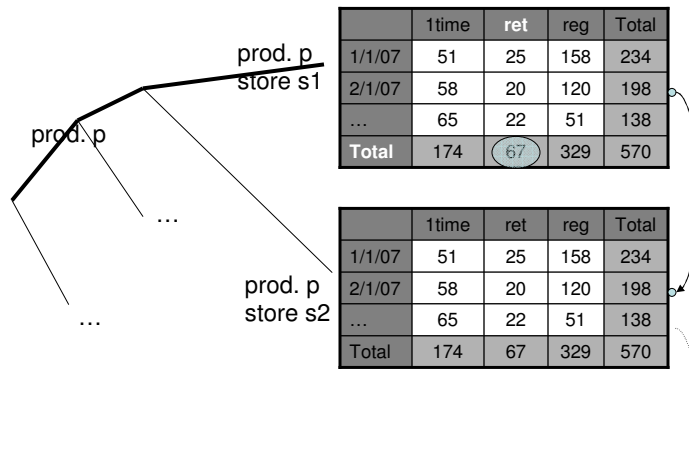


Queries

- Efficiency depends on:
 - does index on sparse dimensions fit into memory?
 - Type of queries:
 - Restrictions on all dimensions
 - Restrictions only on dense
 - Restrictions only on some sparse and dense

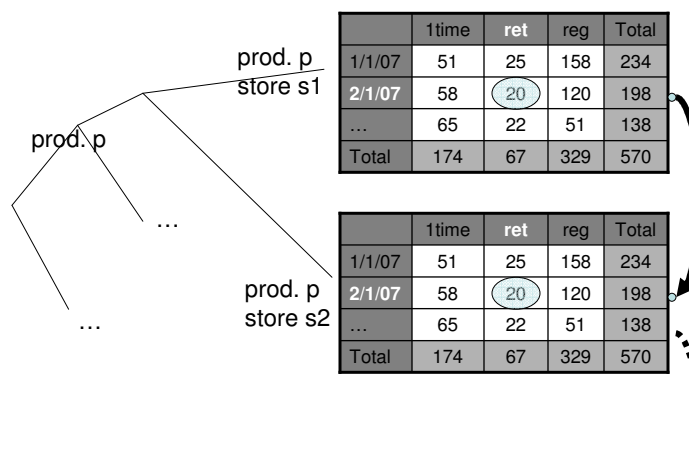
Queries

- Selection on all attributes: (p,s1,ret,all)



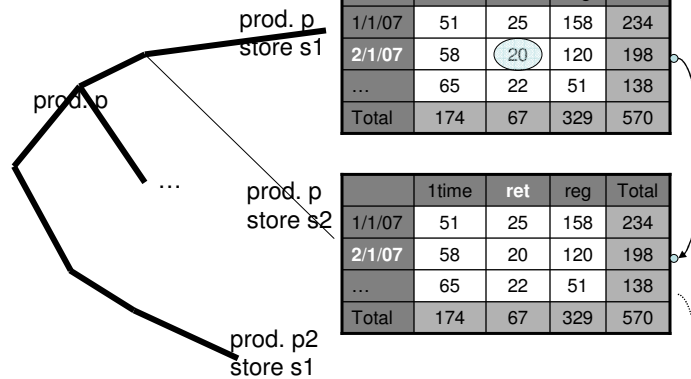
Queries

- Only on dense attributes: (-,-,ret,"2/1/07")



Queries

- Only some sparse and dense attributes:
(-,s1,ret,"2/1/07")

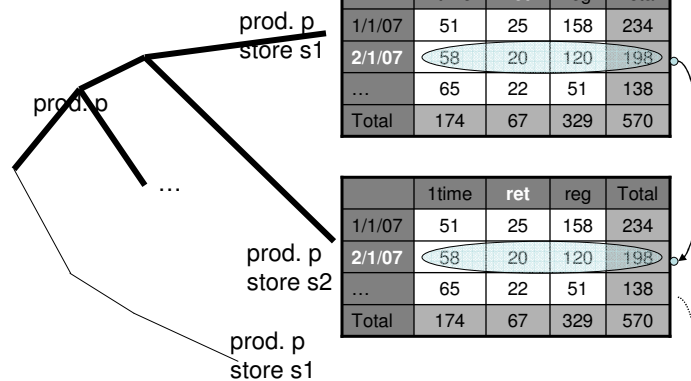


	1time	ret	reg	Total
1/1/07	51	25	158	234
2/1/07	58	20	120	198
...	65	22	51	138
Total	174	67	329	570

	1time	ret	reg	Total
1/1/07	51	25	158	234
2/1/07	58	20	120	198
...	65	22	51	138
Total	174	67	329	570

Queries

- Only some sparse and dense attributes: (p,-,"2/1/07")



	1time	ret	reg	Total
1/1/07	51	25	158	234
2/1/07	58	20	120	198
...	65	22	51	138
Total	174	67	329	570

	1time	ret	reg	Total
1/1/07	51	25	158	234
2/1/07	58	20	120	198
...	65	22	51	138
Total	174	67	329	570

Storing the Cube

- Dense combinations of dimensions can be stored in multi-dimensional arrays
- For every combination of sparse dimensions
 - one sub-cube
- Sub-cubes indexed by sparse dimensions
 - E.g., B-tree
 - Order of the dimensions plays a role

This Lecture

- How is the data stored?
 - relational database (ROLAP)
 - Multi-dimensional structure (MOLAP)
- How can we speed up computation?
 - Indexing structures
 - bitmap index
 - join index

Specialized Indexing Structures

- B-trees, (covered in other courses)
- Bitmapped indices,
- Join indices,
- Spatial data structures (covered later)

Index Structures

- Indexing principle:
 - mapping key values to records for associative direct access
- Most popular indexing techniques in relational database: B+-trees
- For multi-dimensional data, a large number of indexing techniques have been developed: R-trees

Bitmap Indexes

- **Bitmap index:** indexing technique that has attracted attention in multi-dimensional DB implementation table

Customer	City	Car
c1	Detroit	Ford
c2	Chicago	Honda
c3	Detroit	Honda
c4	Poznan	Ford
c5	Paris	BMW
c6	Paris	Nissan

Bitmap Indexes

- The index consists of bitmaps:

ec1	Chicago	Detroit	Paris	Poznan
1	0	1	0	0
2	1	0	0	0
3	0	1	0	0
4	0	0	0	1
5	0	0	1	0
6	0	0	1	0

↑
bitmaps

ec1	BMW	Ford	Honda	Nissan
1	0	1	0	0
2	1	0	1	0
3	0	0	1	0
4	0	1	0	0
5	1	0	0	0
6	0	0	0	1


↑
bitmaps

- Index on a particular column
- Index consists of a number of bit vectors - bitmaps
- Each value in the indexed column has a bit vector (bitmaps)
- The length of the bit vector is the number of records in the base table
- The i -th bit is set if the i -th row of the base table has the value for the indexed column

Bitmap Indexes

- Index on a particular column
- Index consists of a number of bit vectors - bitmaps
- Each value in the indexed column has a bit vector (bitmaps)
- The length of the bit vector is the number of records in the base table
- The i -th bit is set if the i -th row of the base table has the value for the indexed column

Bitmap Indexes



id	name	age
1	joe	20
2	fred	20
3	sally	21
4	nancy	20
5	tom	20
6	pat	25
7	dave	21
8	jeff	26

⋮
data
records

Query:
Get people with age =
20 and name = "fred"

List for age = 20:
1101100000

List for name = "fred":
0100000001

Answer is intersection:
0100000000

Suited well *for domains*
with small cardinality

Bitmap Index

- Size of bitmaps can be further reduced
 - use run-length encoding

1111000111100000001111000 is encoded as
4x1;3x0;4x1;7x0;4x1;3x0

- Can reduce the storage space significantly
 - Logical operations can work directly on the encoding

Bitmap Index – Summary

- With efficient hardware support for bitmap operations (AND, OR, XOR, NOT), bitmap index offers better access methods for certain queries
 - e.g., selection on two attributes
- Some commercial products have implemented bitmap index
- Works poorly for high cardinality domains since the number of bitmaps increases
- Difficult to maintain - need reorganization when relation sizes change (new bitmaps)

This Lecture

- How is the data stored?
 - relational database (ROLAP)
 - Specialized structures (MOLAP)
- How can we speed up computation?
 - Indexing structures
 - bitmap index
 - join index

Join Indexes

- **Traditional indexes:** value \rightarrow rids.
Join indices: tuples in the join \rightarrow to rids in the source tables.
- Data warehouse:
 - values of dimensions of star schema \rightarrow rows in fact table.
- Join indexes can span multiple dimensions

Join

- “Combine” SALE, PRODUCT relations
- In SQL: SELECT * FROM SALE, PRODUCT

sale	prodid	storeld	date	amt
	p1	c1	1	12
	p2	c1	1	11
	p1	c3	1	50
	p2	c2	1	8
	p1	c1	2	44
	p1	c2	2	4

product	id	name	price
	p1	bolt	10
	p2	nut	5

joinTb	prodid	name	price	storeld	date	amt
	p1	bolt	10	c1	1	12
	p2	nut	5	c1	1	11
	p1	bolt	10	c3	1	50
	p2	nut	5	c2	1	8
	p1	bolt	10	c1	2	44
	p1	bolt	10	c2	2	4

Join Indexes

join index

product	id	name	price	jIndex
	p1	bolt	10	r1,r3,r5,r6
	p2	nut	5	r2,r4

sale	rld	prodid	storeld	date	amt
	r1	p1	c1	1	12
	r2	p2	c1	1	11
	r3	p1	c3	1	50
	r4	p2	c2	1	8
	r5	p1	c1	2	44
	r6	p1	c2	2	4

OLAP - Summary

- Data warehouse is a specialized database to support analytical queries = OLAP queries
- Data cube as conceptual model
- Implementation of Data Cube
 - View selection problem
 - Explosion problem
 - ROLAP vs. MOLAP
 - Indexing structures